

Data Structures and Algorithms

Lecture 12: **B⁺ Trees I**

Department of Computer Science & Technology
United International College

Agenda


- Why B⁺ Trees?
- B⁺ Tree Introduction
- B⁺ Tree Operations
 - Search
 - Insertion
 - Deletion

Motivation

- AVL tree with N nodes is an excellent data structure for searching, indexing, etc.
 - The Big-Oh analysis shows most operations finishes within $O(\log N)$ time
- The theoretical conclusion works as long as the entire structure can fit into the main memory
- When the data size is too large and has to reside on disk, the performance of AVL tree may deteriorate rapidly

A Practical Example

- A 500-MIPS machine, with 7200 RPM hard disk
 - 500 million instruction executions, and approximately 120 disk accesses each second
- The machine is shared by 20 users
 - Thus for each user, can handle $120/20=6$ disk access/sec
- A database with 10,000,000 items,
 - 256 bytes/item (assume it doesn't fit in main memory)
 - The typical searching time for one user
 - A successful search need $\log_{\text{base } 2} 10,000,000 = 24$ disk access,
 - Takes around $24/6=4$ sec.
 - This is way too slow!!
- We want to reduce the number of disk access to a very small constant



From Binary to M-ary

- Idea: allow a node in a tree to have many children
 - Less disk access = smaller tree height = more branching
- As branching increases, the depth decreases
- An M-ary tree allows M-way branching
 - Each internal node has at most M children
- A complete M-ary tree has height that is roughly $\log_M N$ instead of $\log_2 N$
 - if $M = 20$, then $\log_{20} 2^{20} < 5$
 - Thus, we can speedup the search significantly

M-ary Search Tree

- Binary search tree has one key to decide which of the two branches to take
- M-ary search tree needs $M-1$ keys to decide which branch to take
- M-ary search tree should be balanced in some way too
 - We don't want an M-ary search tree to degenerate to a linked list, or even a binary search tree
 - Thus, require that each node is at least $\frac{1}{2}$ full!

B⁺ Tree

- A B⁺-tree of order M ($M > 3$) is an M-ary tree with the following properties:
 1. The data items are stored in leaves
 2. The root is either a leaf or has between two and M children
 3. The non-leaf nodes store up to M-1 keys to guide the searching; key i represents the smallest key in subtree $i+1$
 4. All non-leaf nodes (except the root) have between $\lceil M/2 \rceil$ and M children
 5. All leaves are at the same depth and have between $\lceil L/2 \rceil$ and L data items, for some L (usually $L \ll M$, but we will assume $M=L$ in most examples)

There are various definitions of B-trees, but mostly in minor ways.
The above definition is one of the popular forms.

Keys in Internal Nodes

- Which keys are stored at the internal nodes?
 - There are several ways to do it. Different books adopt different conventions.
- We will adopt the following convention:
 - key i in an internal node is the smallest key in its $i+1$ subtree (i.e. right subtree of key i)
- Even following this convention, there is no unique B⁺-tree for the same set of records.

B⁺ Tree Example 1 (M=L=5)

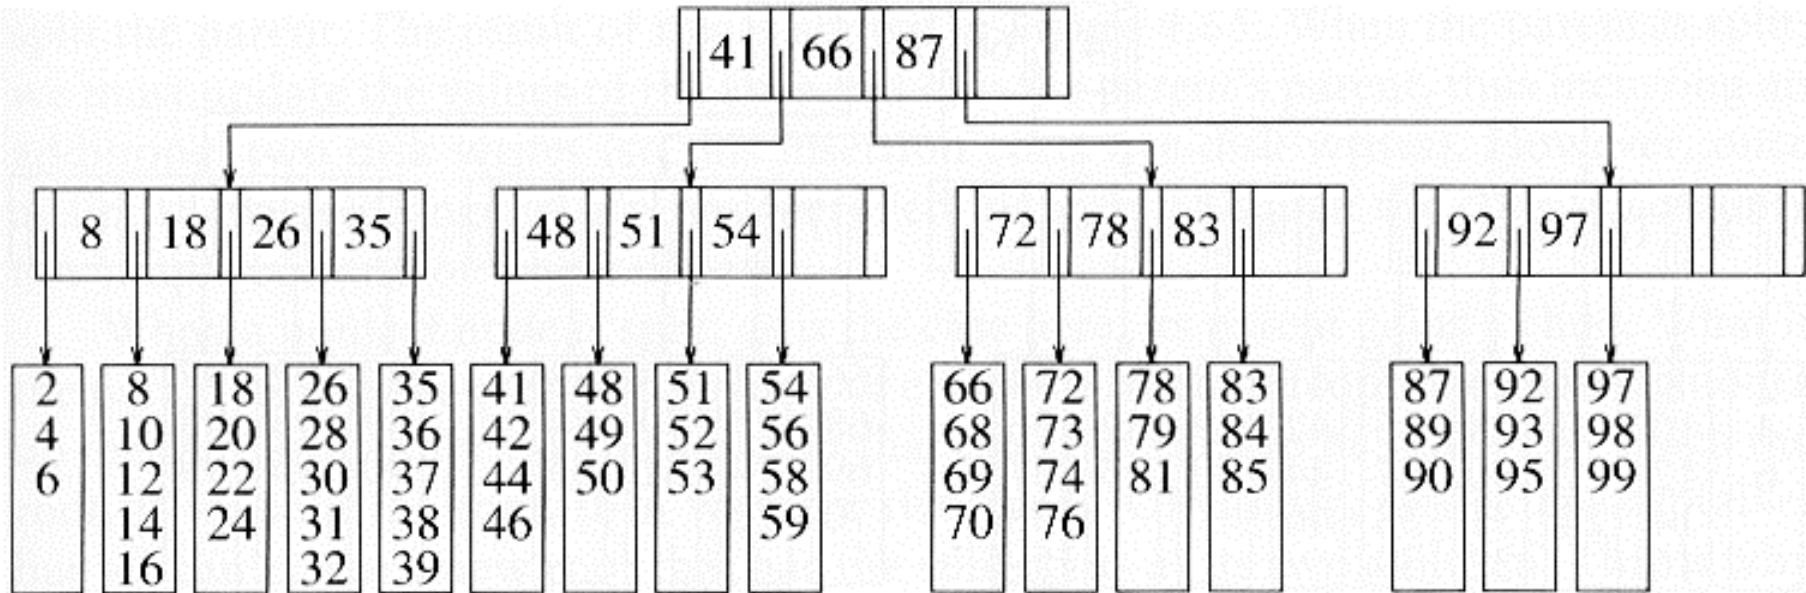



Figure 4.62 B-tree of order 5

- Records are stored at the leaves (we only show the keys here)
- Since $L=5$, each leaf has between 3 and 5 data items
- Since $M=5$, each nonleaf nodes has between 3 to 5 children
- Requiring nodes to be **half full** guarantees that the B+ tree does not degenerate into a simple binary tree

B⁺ Tree Example 2 (M=L=4)



- We can still talk about **left** and **right child pointers**
- E.g. the left child pointer of N is the same as the right child pointer of J
- We can also talk about the **left subtree** and **right subtree** of a key in internal nodes

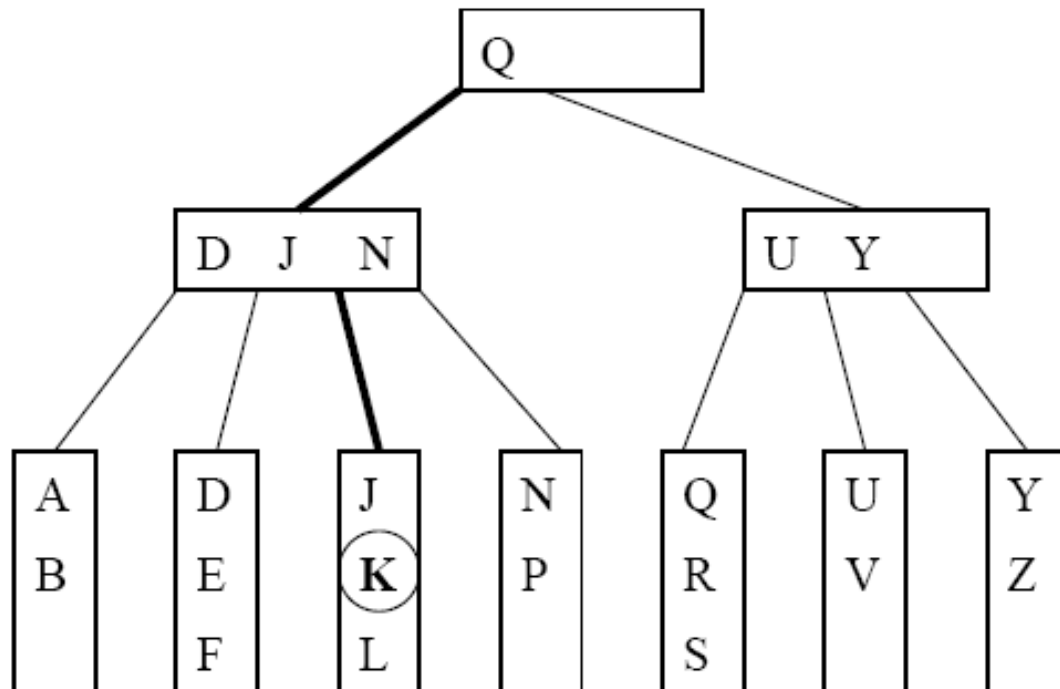
B+ Tree in Practical Usage




- Each internal node/leaf is designed to fit into one I/O block of data. An I/O block usually can hold quite a lot of data. Hence, an internal node can keep a lot of keys, i.e., large M . This implies that the tree has **only a few levels** and **only a few disk accesses** can accomplish a search, insertion, or deletion.
- B⁺-tree is a popular structure used in commercial databases. To further speed up the search, **the first one or two levels of the B⁺-tree are usually kept in main memory**.
- The disadvantage of B⁺-tree is that most nodes will have less than $M-1$ keys most of the time. This could lead to severe **space wastage**. Thus, it is not a good dictionary structure for data in main memory.
- The textbook calls the tree B-tree instead of B⁺-tree. In some other textbooks, B-tree refers to the variant where the actual records are kept at internal nodes as well as the leaves. Such a scheme is not practical. Keeping actual records at the internal nodes will limit the number of keys stored there, and thus increasing the number of tree levels.

Searching Example

- Suppose that we want to search for the key K. The path traversed is shown in bold



Searching Algorithm



- Let x be the input search key.
- Start the searching at the root
- If we encounter an internal node v , search (linear search or binary search) for x among the keys stored at v
 - If $x < K_{\min}$ at v , follow the left child pointer of K_{\min}
 - If $K_i \leq x < K_{i+1}$ for two consecutive keys K_i and K_{i+1} at v , follow the left child pointer of K_{i+1}
 - If $x \geq K_{\max}$ at v , follow the right child pointer of K_{\max}
- If we encounter a leaf v , we search (linear search or binary search) for x among the keys stored at v . If found, we return the entire record; otherwise, report not found.

Insertion Procedure


- Suppose that we want to insert a key K and its associated record.
- Search for the key K using the search procedure
- This will bring us to a leaf x
- Insert K into x
 - Splitting (instead of rotations in AVL trees) of nodes is used to maintain properties of B⁺-trees [next slide]

Insertion into a Leaf




- If leaf x contains $< L$ keys, then insert K into x (at the correct position in node x)
- If x is already full (i.e. containing L keys). Split x
 - Cut x off from its parent
 - Insert K into x , pretending x has space for K . Now x has $L+1$ keys.
 - After inserting K , split x into 2 new leaves x_L and x_R , with x_L containing the $\lfloor (L+1)/2 \rfloor$ smallest keys, and x_R containing the remaining $\lceil (L+1)/2 \rceil$ keys. Let J be the minimum key in x_R
 - Make a copy of J to be the parent of x_L and x_R , and insert the copy together with its child pointers into the old parent of x .

Inserting into a Non-full Leaf (L=3)




Search for O.




Insert O and maintain the order.

Splitting a Leaf: Inserting T




Search for T.




Insert T and the B+ tree condition is violated.


Splitting Example 1



Split the leaf and distribute the keys.



Make S the parent of the two new leaves.




Insert S into the parent.

Maintain the order of keys and child pointers.


Two disk accesses to write the two leaves, one disk access to update the parent

For $L=32$, two leaves with 16 and 17 items are created. We can perform 15 more insertions without another split

Splitting Example 2




Search for M.




Insert M and the B+-tree condition is violated.

Cont'd



Split the leaf and distribute the keys.



Make L the parent of the two new leaves.

However, we cannot just insert L into the parent as it is already full.


=> Need to split the internal node

Splitting an Internal Node


To insert a **key K** into a full **internal node x**:

- Cut x off from its parent
- Insert K and its left and right child pointers into x, pretending there is space. Now x has M keys.
- **Split x** into 2 new internal nodes x_L and x_R , with x_L containing the $(\lceil M/2 \rceil - 1)$ **smallest keys**, and x_R containing the $\lfloor M/2 \rfloor$ **largest keys**. **Note that the $\lceil M/2 \rceil$ th key J is not placed in x_L or x_R**
- Make J the parent of x_L and x_R , and insert J together with its child pointers into the old parent of x.

Example: Splitting Internal Node (M=4)



Insert L and its child pointers into the parent.




Split the parent.

The key J becomes the parent of the two new internal nodes.

Insert J into the next parent.

Cont'd



Split the parent.

The key J becomes the parent of the two new internal nodes.

Insert J into the next parent.

Termination

- Splitting will continue as long as we encounter full internal nodes
- If the split internal node x does not have a parent (i.e. x is a root), then create a new root containing the key J and its two children

